

Problem Set - 4

Please read all of the guidelines carefully before submitting the problem set. (Unless specified) each question is **20 points** and there are **100 points** in total.

Due date: Sunday, February 20, 11:59 PM. Late submissions will be accepted with a penalty! (10% reduction per day – no submissions accepted two days after the deadline.)

Guidelines – Before You Start

- 1) **You should complete the problem set on your own.** Discussing ideas is fine; but, sharing answers and sharing code will be considered as plagiarism.
- 2) You will be using the **Python** programming language. You need to write your codes in an empty **.ipynb** file.
- 3) Make sure that you provide many comments to describe your code and the variables that you created.
- 4) Please use **LaTeX** or **MS Word** to submit your written responses (hand-written responses will not be graded).
- 5) For some of the coding exercises, you may need to do a little bit of “**Googling**” or review the documentation.

Deliverables:

- 1) The code of the problem set in **.ipynb** format (one file)
- 2) Short answers written with **LaTeX** or **MS Word** and exported in **.pdf** format (one file)

Questions

For (most of) the questions below, please use the fake news dataset uploaded on *BlackBoard* (called ‘corona_fake.csv’). You can find the file under ‘Data’ tab.

Please include your code also in your .pdf file (in code blocks).

Data Pre-Processing (40 points)

- 1) **[20 points]** Using the `pandas` package for Python, import the **corona_fake.csv** dataset, and do the following:
 - a) **[5 points]** Import the `nltk` package. Check the documentation: <https://www.nltk.org/>
 - b) **[15 points]** Take a look at the **text** column in the dataset, and do the following:
 - i. **[3 points]** Using `nltk.word_tokenize()`, tokenize the text.
 - ii. **[3 points]** Using the POS-tagging feature (`nltk.pos_tag`), POS-tag the tokenized words.
 - iii. **[3 points]** Using `WordNetLemmatizer` (from `nltk.stem` import `WordNetLemmatizer`) lemmatize the pos-tagged words you obtained above. (*Hint*: If there is no available tag, append the token as is; else, use the tag to lemmatize the token)

- iv. **[3 points]** Using the list of stop words that can be imported (`nltk.corpus` `import stopwords`), remove the stopwords in lemmatized text [**Note:** the language needs to be set as 'english'].].
 - v. **[3 points]** Finally, also **remove numbers, words that are shorter than 2 characters, punctuation, links and emojis**. Finally, convert the obtained list of tokenized+tagged+lemmatized+cleaned list of words back into a joined string (joined by space ' ') and add the result as ***text_clean*** column to your dataset.
- 2) **[20 points]** Let's vectorize the data we produced above by using two approaches: Bag of Words (BOW) and TF-IDF; and, at the end, we will make a prediction:
- a. **[5 points]** Read the following page: <https://en.wikipedia.org/wiki/N-gram>. Explain what an 'n-gram' is and why it is helpful in max. 200 words.
 - b. **[5 points]** Import `CountVectorizer` and `TfidfVectorizer`:

```
from sklearn.feature_extraction.text import
CountVectorizer, TfidfVectorizer
```
 - c. **[5 points]** Using `CountVectorizer`, create three vectorized representations of ***text_clean*** [**set lowercase=True**]:
 - i. One vectorized representation where `ngram_range = (1,1)`
 - ii. One vectorized representation where `ngram_range = (1,2)`
 - iii. One vectorized representation where `ngram_range = (1,3)`
 - d. **[5 points]** Using `TfidfVectorizer`, create three vectorized representations of ***text_clean*** [**set lowercase=True**]:
 - i. One vectorized representation where `ngram_range = (1,1)`
 - ii. One vectorized representation where `ngram_range = (1,2)`
 - iii. One vectorized representation where `ngram_range = (1,3)`

Prediction (20 points)

- 3) **[20 points]** Now, let's use `sklearn.linear_model.LogisticRegressionCV` to do some predictions. Set `cv = 5`, `random_state = 265`, and `max_iter = 1000`, and `n_jobs = -1` (other parameters should be left as default) [**Note:** training size is 70%, test size is 30%, split by `random_state = 265`].
- a. **[10 points]** By using the **three (3)** different versions of the `CountVectorizer` dataset you created above, run logistic regression to predict class labels (fake, true). Report **three (3)** accuracy values associated with each of the regressions.
 - b. **[10 points]** By using the **three (3)** different versions of the `TfidfVectorizer` dataset you created above, run logistic regression to predict class labels (fake, true). Report **three (3)** accuracy values associated with each of the regressions.
 - c. Combine and report all accuracy values in a table (6 values in total).

Theoretical question (40 points)

- 4) **[40 points]** Check the optimizer (solver) functions used by `sklearn.linear_model.LogisticRegressionCV`. For each function, explain in around 100 words what they mean; specifically:
- [8 points]** What does ***newton-cg*** mean?
 - [8 points]** What does ***lbfgs*** mean?
 - [8 points]** What does ***liblinear*** mean?
 - [8 points]** What does ***sag*** mean?
 - [8 points]** What does ***saga*** mean?

Note: For this question you might need to do some online research. It is your job to find out how they work. You are also welcome to use formulas / matrices in your description.