

Introduction to Statistical Machine Learning

CSC/DSCC 265/465

Lecture 8: Supervised Learning – Part V

Cantay Caliskan



Notes and updates

Notes and updates

- Average for **PS1** is **92.37**.
 - Names of graders posted on BlackBoard (under '*List of Graders*')
- Let me know if you are looking for additional resources
- If you have grading-related questions:
 - List of graders for **PS1** have been shared through an announcement
 - Please:
 - Contact your grader first
 - If more clarification is needed, contact our head TA (**Jawahar**)
 - No need to CC me for grading-related questions

- New deadline for the 3rd Problem Set is **Sunday, February 13, 11:59 PM**

Plan for today

- ***Bias and Variance***
- ***Cross-Validation***
- ***Multinomial Logistic Regression***

Plan for today

- ***Bias and Variance***
- *Cross-Validation*
- *Multinomial Logistic Regression*

Bias and Variance

Bias and Variance

- Goal: Hit the target!




Fig. 1: Graphical Illustration of bias-variance trade-off , Source: Scott Fortmann-Roe., Understanding Bias-Variance Trade-off

Bias and Variance

- Idea: Understanding how *bias* and *variance* happen can help us manage reducing them
- **Error due to bias**: Calculated as the difference between the expected prediction of our model and the correct value (=ground truth)
- **Error due to variance**: The variance shows how much the predictions for a given point vary between different realizations of the model
- Question: When is bias more likely? When is variance more likely to happen?

Bias and Variance

- There is a trade-off:
- **Less complex** models (=models with **fewer parameters**) have high bias and low variance
 - Not accurate but *generalizable*
- **More complex** models (=models with **more parameters**) have low bias and high variance
 - *Accurate* but not generalizable
- Question: Where is the optimal model?
- Answer: Somewhere in between

Bias and Variance: Which is worse?

- Most data scientists would try to lower bias at the expense of variance
 - Idea: If I am making the prediction **sometimes** correctly, this is better than getting it wrong **all the time**
- Answer: **Yes** and **No**
 - You can only understand the effects of bias and variance if you realize a model a lot of times
 - In fact: Most people realize a model only once
 - Also: Think about the application
- If you only realize the model once: Long run averages are irrelevant. Bias and variance are equally important.

Bias and Variance

- Ways to deal with **variance**
 - Running several similar models at once / in parallel
 - Examples: Ensemble methods, Bagging, Random Forest
 - Apply **regularization**
- Ways to deal with **bias**
 - Increasing complexity
 - Prefer overfitting over underfitting
 - Known as **model selection**

Plan for today

- *Bias and Variance*
- ***Cross-Validation***
- *Multinomial Logistic Regression*

Cross-Validation

Cross-Validation

- Definition: Different model validation techniques for assessing how the results of a statistical analysis (model) will generalize to an independent data set
- Important: It is not a tool for estimating model coefficients
- Usually used in the context of ***prediction***
- Why is it helpful? What is the goal?
 - Helps us to evaluate the quality of the model
 - Helps us to select the model which will perform best on unseen data
 - Helps us to avoid ***overfitting*** and ***underfitting***
 - Helps us to have a model that is low on ***bias*** and ***variance***

Cross-Validation

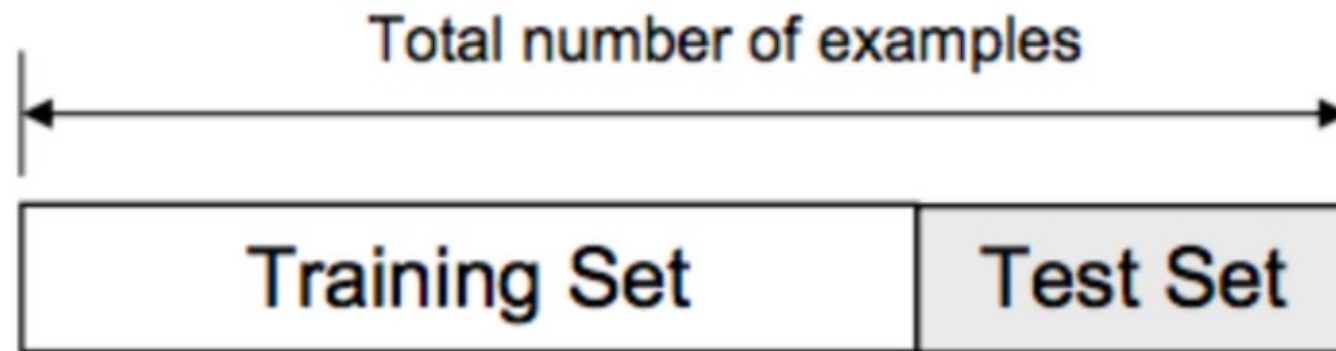
- Idea:
 - You have a ***sample dataset***
 - This dataset represents the characteristics in the ***whole population***
 - This sample dataset may have some "random" differences from the population data
 - Goal: To create an explanatory / predictive algorithm from the sample dataset
 - The algorithm will almost always perform less well when applied on the whole population
 - Why? Because of random error ...
 - You do cross-validation -> To check the ***amount of error*** resulting from the random errors

Cross-Validation Strategies

- 1. Train / Test split**
- 2. K-fold cross-validation**
- 3. Leave one group out**
- 4. Stratification**

Train / Test Split

- **Train / Test Split**
 - Number of groups: 2
 - We split the data into two sets: Training and test data
 - Samples between train and test data sets do not overlap
 - Important not to have duplicate samples in our dataset



Train / Test Split

- **Train / Test Split** (or Holdout)
 - Disadvantage:
 - What if the train / test split isn't random
 - This will result in ***overfitting***
- A good choice, if we have enough data
- Implementation in **Python**:
`sklearn.model_selection.train_test_split`

Pseudocode


- How should we go forward?
 - 1) **Split** the data into training test and validation datasets
 - 2) Find a set of **parameters** that you think would work well with the model
 - 3) Run the model by using a **cost function**
 - 4) Check the cost value obtained from the **test dataset**
 - 5) Report the cost value

K-fold cross-validation (#groups = k)

- Removing a part of the dataset for validation means:
 - Risk of underestimating the amount of ***overfitting***
- Reducing the size of the training data means:
 - Risk of overestimating the amount of ***underfitting***
- Solution: K-fold cross-validation
 - Provides ample data for training the model
 - Leaves a lot of data for validation
- Idea: repeated holdout, and average scores after **K** different holdouts

K-fold cross-validation (#groups = k)

- A good choice when we have:
 - Low amount of data / small data set
 - When the choice of optimal parameters are greatly different between folds




- As a general rule, we choose **k=5** or **k=10** ...
- Implementation in **Python**:
`sklearn.model_selection.KFold`

Pseudocode (when we have #k folds)

- What is difference here?
- 1) **Split** the data into training and test datasets
- 2) Find a set of **parameters** that you think would work well with the model
- 3) Place the parameters in a loop structure
- 4) Place the folds in a nested loop structure. Create training and validation datasets from each fold
- 5) Fit the model in each smaller **training dataset**
- 6) Report the **cost** for the test dataset for **each fold**
- 7) Report the **average cost** at the end


Leave one out ($\#groups = len(train)$) - LOOCV

- It is a special case of K -fold when K is equal to the number of samples in our data set
 - We will iterate through every sample in our dataset each time each time using $k-1$ object as train samples and 1 object as test set



Leave one out ($\#groups = \text{len}(\text{train})$) - LOOCV

- Useful, when:
 - We have too little data
 - And the model is fast enough to retrain
- Implementation in **Python**:
 - `sklearn.model_selection.LeaveOneOut`

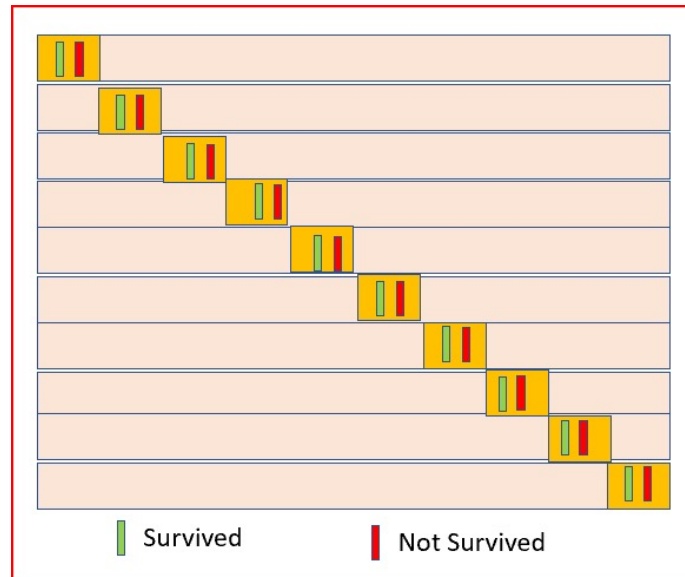


Stratification / Stratified Cross-Validation

- Also called “***Stratified k-fold Cross-Validation***”
- Reminder: ***Stratification*** OR ***stratified sampling***
 - Process of dividing members of population into homogenous subgroups before sampling
- Idea: We re-arrange the data in a way that each fold has a good representation of the whole dataset
 - Stratification forces each fold to have at least some number instances (let's call this ***m***) of each class
 - Ensures that one class of data is not overrepresented especially when the target variable is unbalanced

Example: Stratification

- In a binary classification problem, we want to make sure that each fold has enough number of observations from each class
- Example: Titanic Survivors (!)
 - Survived or not survived
- In a potential ***k-fold*** stratification, each fold needs to have passengers of both sorts



Time series cross-validation

- Splitting time series data randomly is a bad idea:
 - Because of the time dimension
- Solution: Use each group of data created at time t as the training set for each group of data created at time $t+1$

D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9

Time series cross-validation

- We start using smaller training datasets, and make them grow as we go forward ...
- This ensures that we consider the time series aspect of the data for prediction

D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9
D1	D2	D3	D4	D5	D6	D7	D8	D9

Which kind of cross validation?

	Downside	Upside
Test-set	may give unreliable estimate of future performance	cheap
Leave-one-out	expensive	doesn't waste data
10-fold	wastes 10% of the data, 10 times more expensive than test set	only wastes 10%, only 10 times more expensive instead of n times
3-fold	wastes more data than 10-fold, more expensive than test set	slightly better than test-set
N-fold	Identical to Leave-one-out	

Training, Test, Validation Split Ratio

- The split ratio depends on a few factors:
 - The total number of samples in your data
 - The actual model you are training
- Usual practice:
 - Split by **70% by 30%** or **80% by 20%**
 - The larger set is the training set
 - And then (you may) choose a validation set from the training set

Evaluating the model performance

- **Mean Squared Error (MSE)**
- **Root Mean Squared Error (RMSE)**
- **Mean Absolute Percentage Error (MAPE)**
- These are all **cost** functions 😊

Evaluating the accuracy

- We need to understand how well the models give the correct classification
 - And then measure the value of prediction
- This brings us to ***confusion matrix*** ...
- Confusion matrix visualizes the performance of an algorithm by showing true positives, false positives, true negatives, and false negatives

Confusion matrix

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

$$Sensitivity = \frac{TP}{TP + FN}$$

$$Specificity = \frac{TN}{TN + FP}$$

		Actual Values	
		Positive (1)	Negative (0)
Predicted Values	Positive (1)	TP	FP
	Negative (0)	FN	TN

- Goal (usually): Look for a balance between sensitivity and specificity

Confusion matrix

Confusion Matrix		Target			
		1	0		
Logit	1	261	64	Positive Predictive Value	0.803076923
	0	81	485	Negative Predictive Value	0.856890459
		Sensitivity	Specificity	Accuracy=	
		0.763157895	0.88342441	0.837261504	

Titanic Survivors

- Harder to be accurate with higher number of classes
- Goal: Always look for a balance in misclassifications

		Predicted Labels									
		0	1	2	3	4	5	6	7	8	9
True Labels	0	987	1	2	0	0	0	2	0	7	1
	1	0	977	7	2	3	2	0	2	6	1
	2	2	3	976	4	4	0	1	4	6	0
	3	0	1	18	951	0	14	0	3	9	4
	4	0	1	2	0	979	0	2	0	3	13
	5	3	0	3	9	5	968	2	0	5	5
	6	1	3	2	0	0	7	982	0	5	0
	7	3	4	3	0	13	0	0	969	0	8
	8	2	6	4	7	3	5	2	3	966	2
	9	1	1	2	6	12	2	0	8	5	963

MNIST Hand-written Digits

Even better: F1

- If we are looking for “some kind of” balance between sensitivity and specificity:
 - We can use the ***F1 score***
 - Also called F-score or F-measure
- Formally: Harmonic mean of ***precision*** and ***recall***

$$F1 = \frac{2 TP}{2 TP + FP + FN}$$

Reminders

- Please work on the problem set
- Read ***Chapter 12***